



Knative 101: What it is, and what it will be

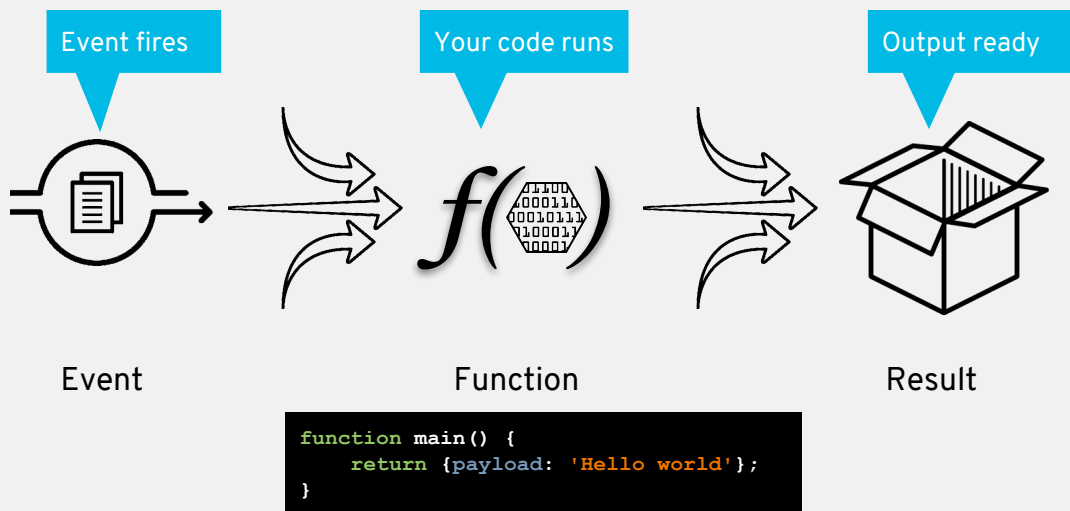
Giuseppe Bonocore
Solution Architect

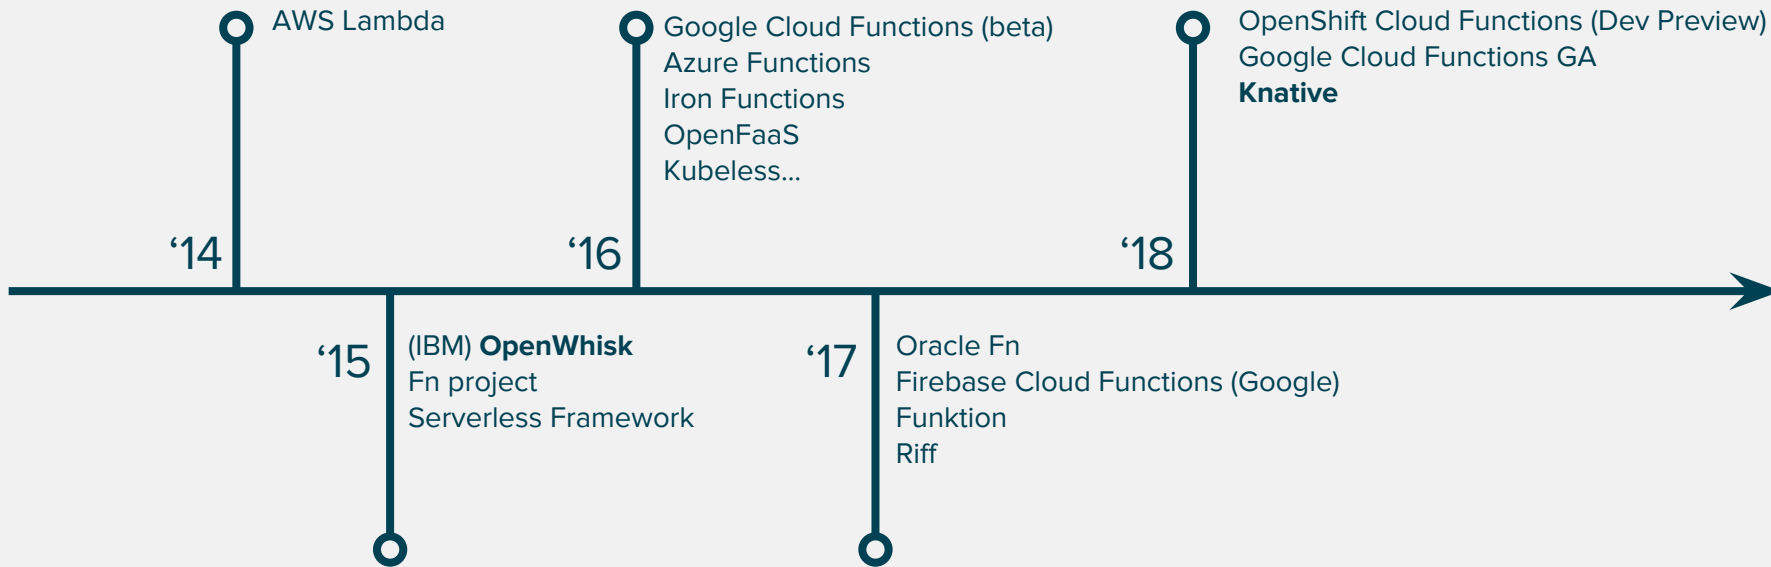


A wide-angle photograph of a large, empty server room. The room features a high ceiling with a grid of square tiles and several rectangular fluorescent light fixtures. The floor is also covered in square tiles, with some areas showing raised access floor grates. On the left side, there is a long row of server racks, some of which are labeled with numbers like 93 and 91. A green exit sign is visible above a door on the left. In the center of the room, there are several white support poles. On the right side, more server racks are visible, also labeled with numbers like 83, 85, 87, 89, 91, and 93. The overall atmosphere is clean and modern.

A Serverless Datacenter !

How does it work ?





Is Serverless Open source?




CLOUD NATIVE
COMPUTING FOUNDATION

Standards!

Standards!

Standards!



cloudevents



redhat.



THE
LINUX
FOUNDATION

Tools



Security




Framework



Hosted

Installable

Platform



s.cncf.io

Serverless computing refers to a new model of cloud native computing, enabled by architectures that do not require server management to build and run applications. This landscape illustrates a finer-grained deployment model where applications, bundled as one or more functions, are uploaded to a platform and then executed, scaled, and billed in response to the exact demand needed at the moment.

 **CLOUD NATIVE Landscape**

 **CLOUD NATIVE COMPUTING FOUNDATION**

 **Redpoint**

Cloud Native Landscape



Welcome to Knative

Knative (pronounced kay-nay-tiv) **extends Kubernetes** to provide a set of **middleware components** that are essential to build **modern, source-centric, and container-based** applications that can **run anywhere**: on premises, in the cloud, or even in a third-party data center.



redhat®



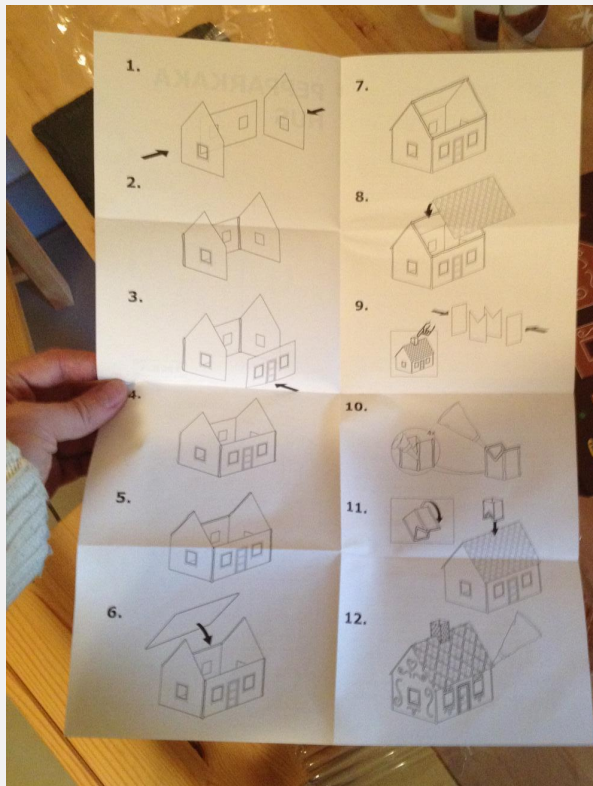
Google

Pivotal



Build

A pluggable model for building artifacts, like jar files, zips or containers from source code.





Serving

An event-driven model that serves the container with your application and can "scale to zero".





Events

Common infrastructure for
consuming and producing
events that will stimulate
applications.



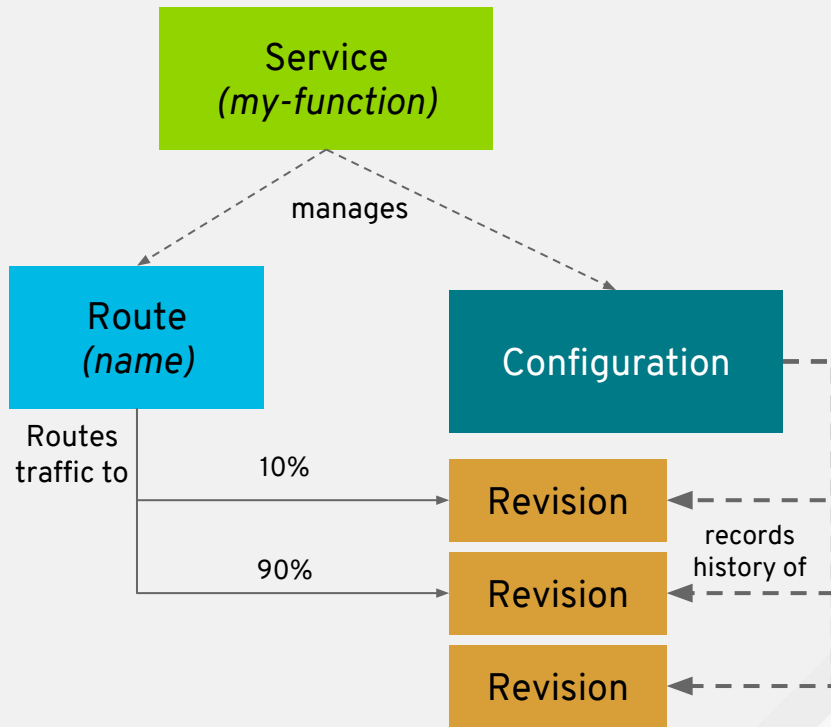
Build

- A **Build** is a list of containers run in-order, with source mounted in
- Implemented as a Kubernetes Custom Resource Definition (**CRD**).
- **BuildTemplates** provide reusable, parameterized recipes that can be used to create Builds
- Pipelines? Maybe
- Aka **S2I** (for the OpenShifters)

```
apiVersion: build.knative.dev/v1alpha1
kind: Build
metadata:
  name: example-build
spec:
  serviceAccountName: build-auth-example
  source:
    git:
      url: https://github.com/example/build-example.git
      revision: master
  steps:
  - name: centos-example
    image: centos
    args: ["centos-build-example", "SECRETS-example.md"]
  steps:
  - image: quay.io/example-builders/build-example
    args: ['echo', 'hello-example', 'build']
```

Serving

- Leverages Istio
- **Configurations** maintains the desired state for your deployment. Modifying a configuration creates a new revision.
- **Revisions** represent an immutable snapshot of code and configuration
- **Routes** configure ingress over a collection of Revisions and/or Configurations
- **Services** (nope, not K8s services) are top-level controllers that manage a set of Routes and Configurations to implement a network service.



Eventing

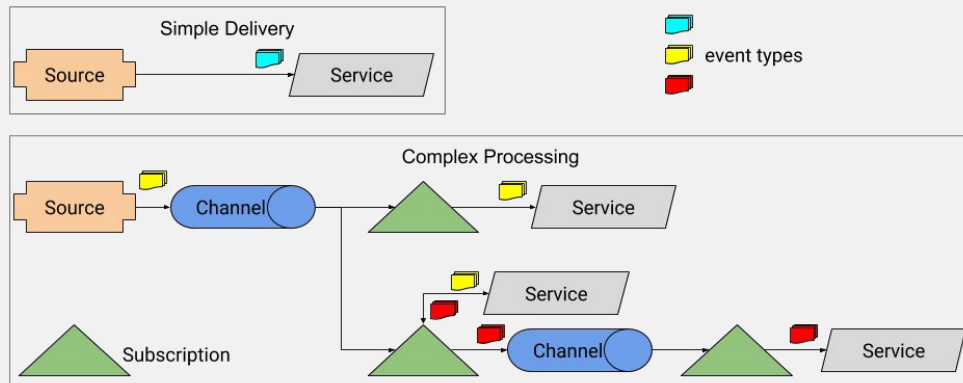
- Design goals consistent to the specification of

CNCF Cloudevents

- **Source:** Produce the events

- **Event Consumers**

- **Addressable:** Receives and ack (K8s services)
- **Callable:** Receives and transform




- **Channel:** named endpoint for event forwarding and persistence layer. Implemented by Kafka, AMQP...
- **Subscription:** Registration between channels and services or other channels



CRDs

- Configuration
- Revision
- Route
- Service
- Build



- III. Config: strict separation of config from code
- I. Codebase: different versions may be active
- VIII. Concurrency: Scale out via the process model
- IX. Disposability: fast startup and graceful shutdown
- V. Build, release, run: Strictly separate build and run

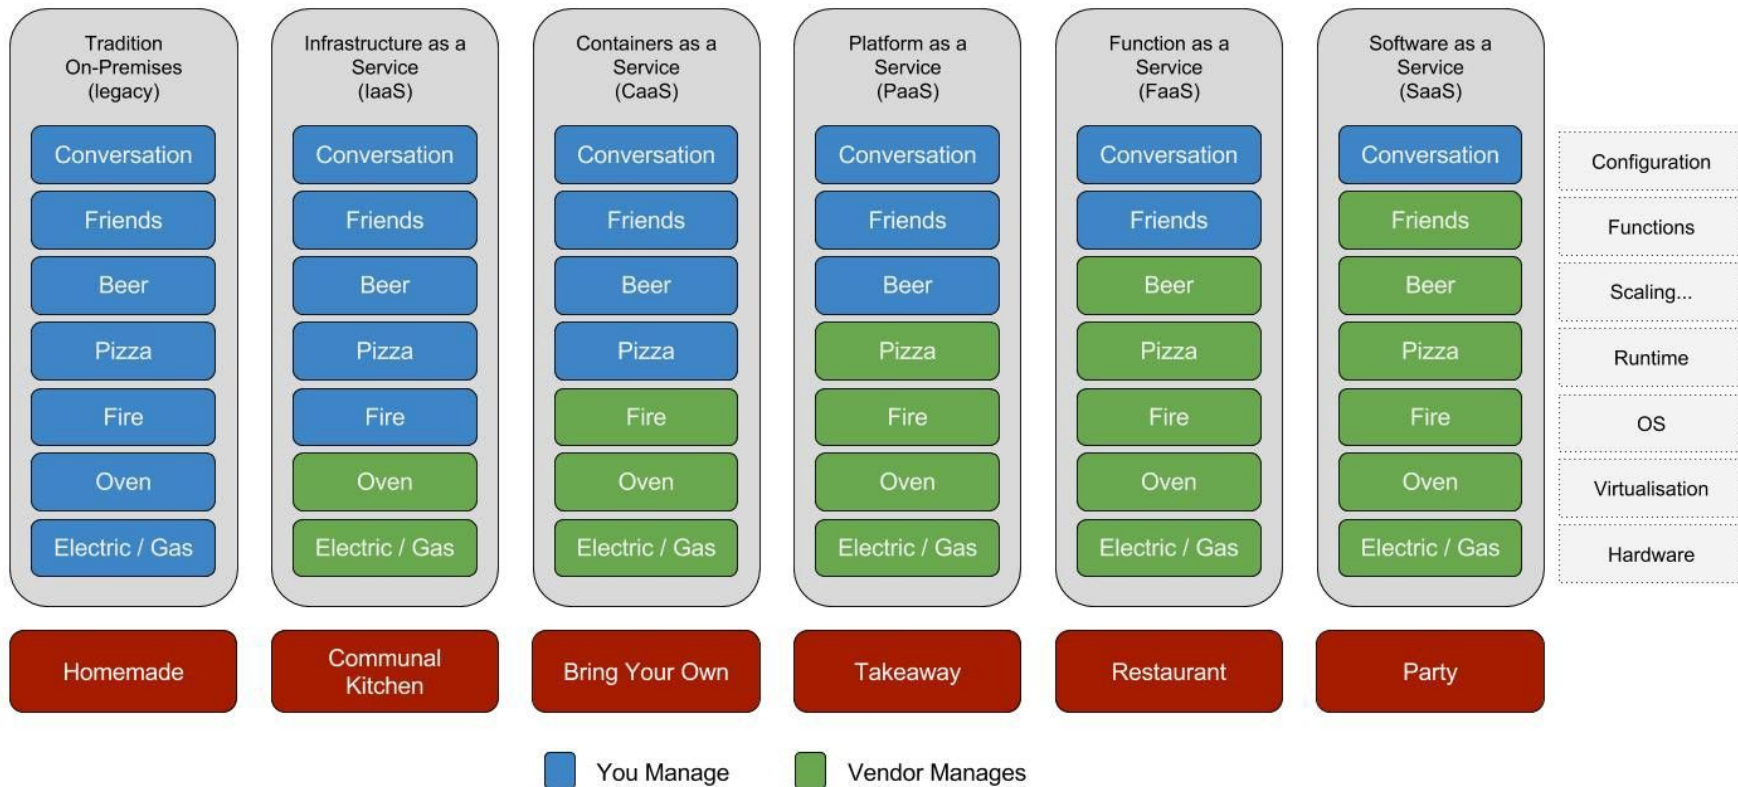


<https://12factor.net>



Pizza as a Service 2.0

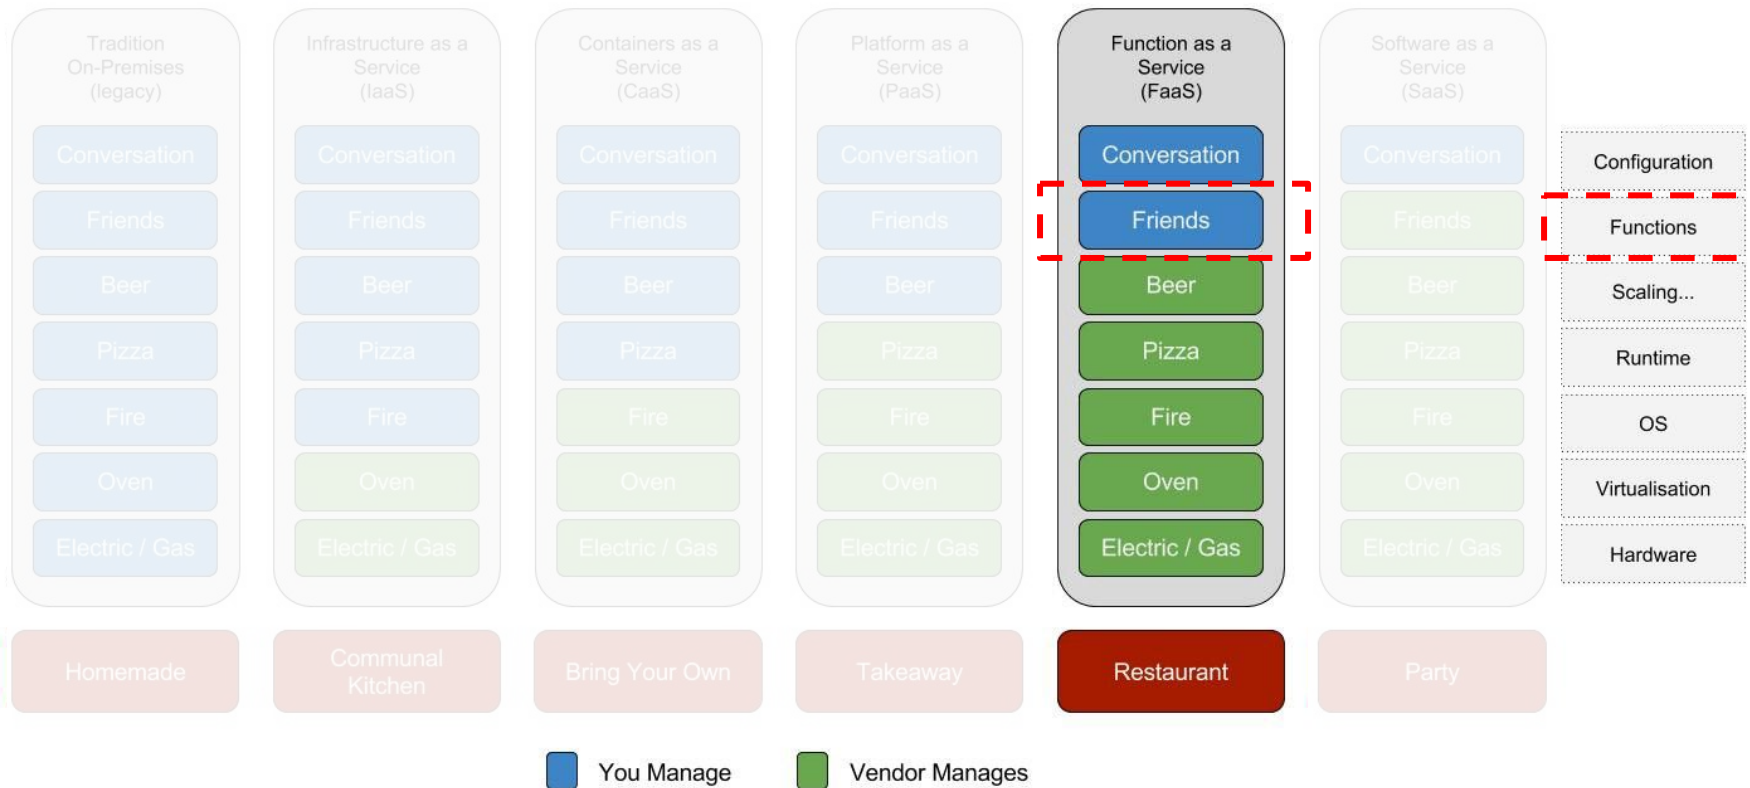
<http://www.paulkerrison.co.uk>





Pizza as a Service 2.0



<http://www.paulkerrison.co.uk>



FaaS

{ f(x) }

Serverless



Services

- User experience
- Services
- No pods for services
- No pods for functions (on user projects)
- Debugging/IDE Integration
- API Gateway Integration
- Billing/Charging model
 - Per function call
 - Per execution time
 - Resource consumption

Function as a Service

Apache
OpenWhisk

Camel K

Riff

Kubeless

OpenFaaS

...

Red Hat OpenShift

Knative

Events

Build

Serving

Automated
Operations

RH MW Services (Operator
backed)

ISV Services (Operator
backed)

Operator Framework

Istio

Hybrid Install / Ops

Install / Upgrade

Network / CNI

Ops & Dev Consoles

Security / Auth

Storage / CSI

Kubernetes

Red Hat Enterprise Linux or Red Hat CoreOS

Common use cases

- Processing **web hooks**
- Scheduled **tasks** (a la cron)
- **Data transformation**
- Mobile **image manipulation**
(compression, conversion, and so on)
- Voice packet to JSON transformation
(Alexa, Cortana, and so on)
- Mobile **video analysis**
(frame-grabbing)
- **PDF generation**
- Mobile/MBaaS /single-page apps
- **Chat bots**

Web

Mobile

IoT

DevOps Automation




**Focus on convenience and
business value, no distractions.**

Asynchronous, concurrent, easy to parallelize into independent units of work

When not to use Serverless

- **Real-time**, ultra-low latency applications
- **Long running tasks** that can't be split into steps
- Advanced or **complex** observability and monitoring requirements
- **Memory or CPU requirements** are very demanding and specific
- Can't deal with cold-start



- bonocore@redhat.com

Q&A

<https://blog.openshift.com/knative-serving-your-serverless-services/>

<https://github.com/knative>

<https://github.com/redhat-developer-demos/knative-tutorial>

*“Serverless computing is a cloud-computing execution model in which **the cloud provider runs the server**, and dynamically manages the allocation of machine resources. **Pricing** is based on the actual amount of resources consumed by an application, rather than on pre-purchased units of capacity. [...]*

*Serverless computing can **simplify** the process of **deploying code** into production. Scaling, capacity planning and maintenance operations may be hidden from the developer or operator.”*

Wikipedia

CNCF Serverless Whitepaper v1.0:

*“**Serverless computing** refers to the concept of **building and running applications that do not require server management**. It describes a finer-grained deployment model where **applications, bundled as one or more functions**, are uploaded to a platform and then **executed, scaled, and billed** in response to the exact demand needed at the moment.”*

